THE INSURABILITY LAYER FOR AGENTIC AI

Make your agents insurable.

Your carriers have filed to exclude AI. RiskMandate measures what your agents can actually do, proves the controls, prices the exposure, and produces the evidence an underwriter will accept.

Design Partner Programme · Cohort One · Five to seven partners · August 2026

The Problem

Cover is being withdrawn. Evidence buys it back.

FILINGS

60+

P&C groups filed to exclude AI

ISO forms CG 40 47, CG 40 48 and CG 35 08 took effect 1 January 2026. Some exclusions on D&O and E&O are absolute.

APPROVALS

80%

of those filings cleared review

Whether AI is covered is now decided policy by policy, jurisdiction by jurisdiction, at each renewal date.

THE GATE

1

question decides the outcome

Can you show what your agents reach, and prove the controls hold? Everything else follows from that answer.

Affirmative cover exists. Every carrier writing it asks the same question — and most teams cannot answer it in writing.

The Difference

Questionnaires describe. Evidence prices.

	Questionnaire	RiskMandate
Where the answer comes from	**Someone's recollection.** A security lead fills in a form once a year, from memory and a spreadsheet last accurate in March.	**Your environment.** Read-only connectors derive what each agent can actually reach, refreshed as permissions change.
What the underwriter gets	**A yes or no.** No severity, no aggregation, no way to tell a contained agent from one holding standing production access.	**A priced exposure.** Loss bands per agent, aggregate across the fleet, and concentration in shared models and vendors.
What happens at claim	**The answers get tested.** Application warranties are examined after the loss. This is where cover is most often lost.	**The record holds.** Timestamped control state, containment tests and an immutable action log, assembled beforehand.

Six levels. One number.

Level	Name	Score range
L0	Unmapped	0–15
L1	Inventoried	16–35
L2	Scoped	36–55
L3	Attested	56–72
L4	Quantified	73–88
L5	Underwritten	89–100

L0–L1
Declination, or an exclusion attaches. Nothing for an underwriter to assess.

L2
Affirmative endorsement at a low sub-limit. First point a broker can market the risk.

L3
Standalone AI liability. Retention drops. Warranties survive a forensic audit.

L4–L5
Priced on your own data, not class rates. Captive-eligible at the top.

What Gets Measured

Five dimensions. Weighted by price impact.

30%

Exposure containment

How bad one agent can get. Reachable actions, systems touched, dollar value of authority, and whether the damage is reversible.

Drives severity

20%

Authority definition

Whether every agent has a written mandate: purpose, permitted actions, data scope, and where a human must intervene.

Drives frequency

20%

Attestation integrity

Whether control state is proven continuously rather than asserted once a year, with logs that hold up under examination.

Drives warranty credibility

20%

Loss quantification

Expected loss per agent as a range, not a point. A wide range is itself a finding — it means clarity is the next thing to fund.

Drives pricing

10%

Accountability

Who owns each accepted risk, for how long, at what retention. Acceptance without a named owner and expiry is not acceptance.

Drives legal standing

Derived from your environment through read-only connectors — never from a self-assessment.

The Engagement

What a design partner gets.

01

Your Insurability Index

A score from 0–100 across all five dimensions, with every number decomposing to the evidence behind it.

02

A measured blast radius

Per agent: reachable actions, systems touched, dollar value of authority, reversibility class, standing versus time-bound privilege.

03

The gap list

Exactly what closing the distance to the next level requires, ordered by what moves a price most.

04

A renewal report

Written for an underwriter rather than a security team. Your broker submits it without rewriting.

05

The acceptance ledger

Every risk accepted by a named owner for a stated interval, with expiry forcing re-decision.

06

Founder-level access

Direct line to the team. Your environment shapes what gets built next.

The Fit

Who this is built for — and who it is not.

A GOOD FIT

- Agents already in production touching money movement, customer data or regulated records
- A renewal inside six months with an AI exclusion attaching or threatened
- Identity and posture tooling already deployed for us to read
- A risk owner outside security — GC, CFO or Head of Risk
- Able to grant anonymised exposure data rights

NOT YET

- Agents still in pilot — there is nothing to measure
- No inventory and no controls deployed. The Index returns near-zero and the gap list becomes a two-year roadmap
- A renewal that just closed, so no forcing function exists
- Agents confined to internal productivity tooling
- A captive that self-insures, so the exclusion does not bite

We would rather tell you it is too early than take the engagement.

How It Works

Four weeks to a number.

Week 1

Connect

Read-only connectors to your identity provider, cloud IAM and agent posture tooling. Setup is typically under a day. We never sit in the request path.

Week 2

Measure

Agent discovery and blast-radius mapping. What each agent can reach, what authority it holds, whether the damage is reversible.

Week 3

Score

Your Insurability Index across all five dimensions, with the derivation shown for every number.

Week 4

Deliver

The gap list and a renewal report your broker can submit. Then quarterly re-assessment as the estate changes.

Deliberately shorter than a renewal window. That is the whole point of the timeline.

The Ask

What we need from you.

Read-only connector access	To your identity provider, cloud IAM and agent posture tooling. Nothing inline, nothing in the request path.
Anonymised exposure data rights	Agent scope and permission metadata — never the contents of what your agents process. This is the one term we hold firm on.
Two hours a month	One working session to review findings and challenge where our score disagrees with your judgement.
Honesty about what breaks	You are shaping the product. Telling us where it is wrong is the point of a design partnership.

In return: the programme is free for cohort one, capped at five to seven partners, with preferential pricing locked for two years afterwards.

Straight Talk

What we are not.

We are not an insurance policy.

We are not a carrier, a broker or an MGA, and we do not sell or place cover. We measure insurability and produce the evidence underwriters price against. Your broker and carrier relationships stay exactly as they are.

We do not replace your security tooling.

We connect to what you already run and translate its output into exposure an underwriter can read. If you have no controls at all, we will tell you that before you buy anything.

The model is not yet actuarial.

Weights are set by underwriting judgement today and re-fit as loss experience accumulates. We say so openly rather than implying a precision that does not yet exist. Design partners are how that changes.

We never sit in the runtime path.

All connectors are read-only and out of band. A governance layer that can take your agents down is a new source of the risk it was bought to measure.

Know your number before your carrier does.

A first assessment returns your Insurability Index, the gap list to the next level, and a renewal report written for an underwriter rather than a security team.

hello@riskmandate.ai

riskmandate.ai

Know the risk. · Name the owner. · Enforce the mandate.